

MemoryStream and Memory-Mapped Files

Mohamad Halabi
Microsoft Integration MVP
@mohamadhalabi



pluralsight 
hardcore dev and IT training

MemoryStream

- **Simple and commonly used**
- **Backing store: in-memory buffer**
 - Very fast: no need for I/O read/write operations
 - Caution: not suitable for large data size
- **Defies an important streaming advantage:**
 - No reading data in chunks (all data is already in memory!)
- **So, should we ever use MemoryStream?**
 - Definitely a good choice for non-persistent small-size data

Usage Scenarios

- **Scenario 1: random access to data chunk from a non-seekable stream**
 - Read data chunk and store in an array
 - Wrap a MemoryStream around the array
 - MemoryStream is seekable
- **Scenario 2: random access to data from I/O call; ex: web service or database**
 - Wrap data in MemoryStream for fast random access

Flush

- Flush has no implementation in MemoryStream

```
/// <summary>  
/// Overrides <see cref="M:System.IO.Stream.Flush"/> so that no action is performed.  
/// </summary>  
/// <filterpriority>2</filterpriority>  
[__DynamicallyInvokable]  
public override void Flush()  
{  
}
```

Why Are Memory-Mapped Files Discussed?

- Memory-mapped files are not stream types
- Provide similar features to files and pipes
 - So memory-mapped files are discussed to understand design alternatives

Memory-Mapped Files

- Types found in `System.IO.MemoryMappedFiles` namespace
- Two key features:

Feature 1

Fast random access to files

`FileStream` allows random access to files

Memory-mapped files vs. `FileStream`

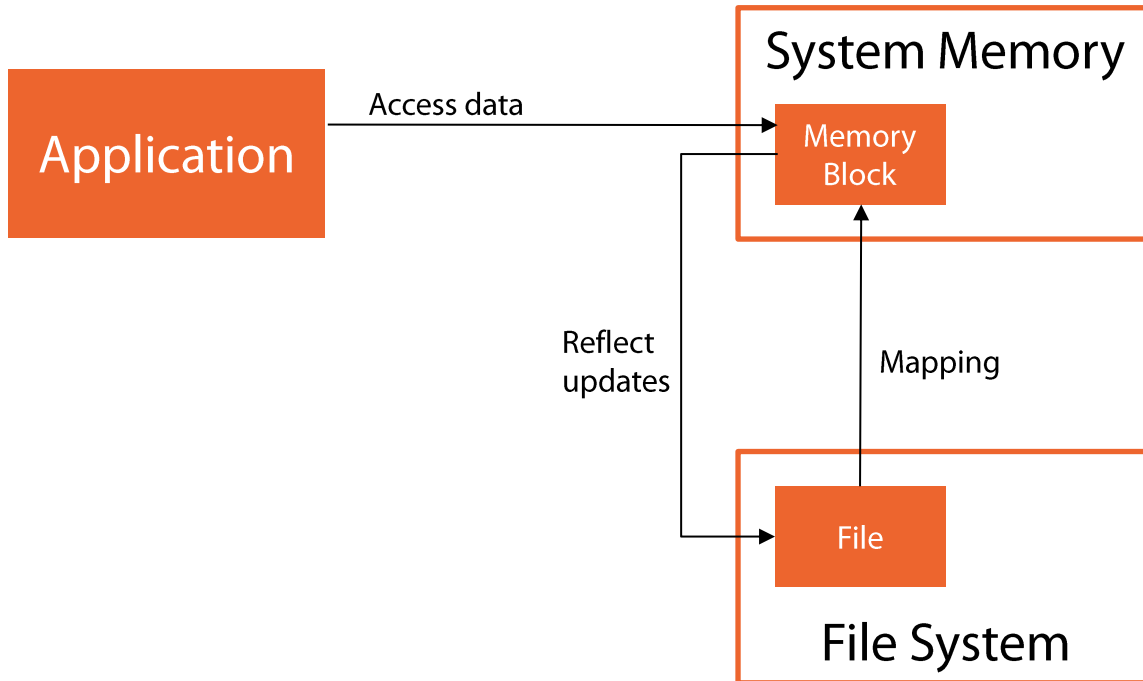
Feature 2

Shared memory between processes on same machine

Pipes provide a shared memory between processes

Memory-mapped files vs. pipes and `PipeStream`

File Access



Memory-Mapped Files vs. FileStream

FileStream

Optimized for sequential file access

I/O required to access data

Not thread-safe because of pointer moves while writing and reading


Memory-Mapped Files

Better performance for random file access

Faster memory access

Data can be read in a multi-threaded fashion

Memory Sharing



- Despite the name memory-mapped "file", there is no system file

Memory-Mapped Files vs. Pipes

Memory-Mapped Files

Allow same-machine communication only

Not stream-based. Shared memory block

Most efficient for single machine communication

Pipes

Allow cross-machine communication

Stream-based

Using the API

- Use MemoryMappedFile's **CreateNew** instead of **CreateFromFile**

```
using (MemoryMappedFile mmf = MemoryMappedFile.CreateNew("MemoryLocation1", 1000))
using (MemoryMappedViewAccessor accessor = mmf.CreateViewAccessor())
{
    accessor.Write(0, 100);
    Console.ReadLine();
}
```

// the second process

```
using (MemoryMappedFile mmf = MemoryMappedFile.OpenExisting("MemoryLocation1"))
using (MemoryMappedViewAccessor accessor = mmf.CreateViewAccessor())
    Console.WriteLine(accessor.ReadInt32(0)); // returns 100
```

Summary

- **MemoryStream is often overused**
 - Suitable for small-size data
 - Not suitable for large-size data
- **Memory-mapped files are not stream types**
 - Fast random file access
 - Fast access to shared memory data between processes
 - Preferred over FileStream and Pipes (in terms of performance)